



The effect of supplementary data on Acoustic Event Classification through Machine Learning

Final Year Project Report - MECH4841 Part B

JUNE 2019

Christopher Caelli¹

¹ *Student of Mechatronics Engineering,
The University of Newcastle, Callaghan, NSW 2308, AUSTRALIA
E-mail: Christopher.Caelli@uon.edu.au*

Mandatory Dot Point Summary

flagforreview [FLAGGED FOR REVIEW]

citationneeded [citation needed]

Tense: Past, Future etc. Search "was", "is" "ed" etc. Check contraction. it's, don't etc. check for our, we, I, they're, your, you, etc. check for st-ray - inca-se the-y block sh-t check ref and replace with fref, tref, aref etc.

As per the FYP Handbook:

I did:

- I learnt and prototyped a machine learning solution that takes audio and physical data (GPS, accelerometer, gyroscope) to compare whether the supplementary data could improve the classification of sound, as measured by the F1 score.
- I learnt and prototyped a pipeline to record data, process it, classify it, and output it.
- I used my mechatronics subject matter learnings to best apply machine learning to this problem. This includes Mechatronics Design (trade off and evaluations from MCHA3000), sound pre-processing, preprocessing Inertial Measuring Unit data, and how Audio waveforms can be treated as energy, and statistics from MECH2450/MCHA3900.
- I learnt and used System Engineering techniques to elicit requirements from my university and work stakeholders

I helped:

- I managed and worked with another employee to implement a proof of concept app to record data as an input into a this project. This involved using my software engineering skills and project management skills to create a schema, develop the android app, and present it to the relevant stakeholders.

Executive summary

The project investigated whether supplementary data (Accelerometer, Gyroscope, GPS) will improve a classifiers ability to classify an "Acoustic Event" in an Acoustic Event Detection / Classification (AED/C) problem. The motivation behind this was to improve AED/C without onerous microphone requirements, to enable more wide spread comercial use of AED/C.

Section 2 of the project begins by evaluating the research problem for a car use case, and eliciting the design requirements through system engineering principles.

Section 3 reviews current research. Presently the field is expanding with broader advances in the wider machine learning community, however AED/C still remains a non-trivial problem that doesn't have a viable solution yet. Comparatively to the wider machine learning community, Acoustic Event Detection is very limited in active researchers. This specific issue hasn't been explicitly answered.

Section 4 details selecting the desired output, the required data input, a metric to score by, a classifier, optimising hyperparameters against the previous options, and finally training and validating against the data. For this project:

- The desired output is a label (or labels) detailing the Acoustic Event in the sample (e.g. The labels found in F like 'Walking').
- The required data is the ExtraSensory dataset, or a custom dataset as recorded by the project's App.
- The score selected is the F1 score. This is due to it being an industry standard.
- The classifier chosen is the MLPClassifier which is a feedforward neural network, trained via backpropagation.
- Hyperparameters of the MLPClassifier are chosen via a Random Search Cross Validation.
- Training and Validation was run with Early Stopping using a 10% validation set, using 59 out of the 60 ExtraSensory users.
- Final testing is done with the 60th user.

Section 5 the desired results are a difference [FLAGGED FOR REVIEW] in the F1 score, after comparing a classifier with and without the supplementary data. This is calculated through analysis of F1 Scores of a binary MLPClassifier, and a multi-label MLPClassifier, both trained on the ExtraSensory dataset using the extact method described in Section 4. The output of the classifiers are also analysed and found to be a valid multilabel classifier for realworld use.

Section 6 evaluates the effect supplementary data had on improving classification. It will discuss how multi-sensor AED/C is effective but requires more datasets to test smaller, more specific tasks. [FLAGGED FOR REVIEW]

Section 7 will present how these results will contribute to the field, and suggested future work into new datasets with a greater number of labels.

Acknowledgements

I'd like to thank my partner Brigid for encouraging and supporting me throughout this project.

I'd also like to thank each and every mechatronics Academic staff member at the University of Newcastle . The teaching staff care about their students, and I want that to grow.

Contents

1	Introduction	7
2	Problem Analysis	7
2.1	Current 18 Words	7
2.2	Tree Diagram	8
2.3	Quad of Aims	8
2.4	Input Output Analysis	8
2.5	Affinity Diagram	9
2.6	Systems Map	9
2.7	Sequence Diagram	10
2.8	N ² Analysis	10
2.9	Spray Diagram	10
2.10	Matrix Diagram	10
2.11	Final Problem Analysis and Discussion	10
3	Review Of Literature	11
3.1	Machine Learning	11
3.2	Acoustic Event Detection and Acoustic Event Classification	12
3.3	DCASE Competitions (Detection and Classification of Acoustic Scenes and Events)	12
3.4	Leading Research, DCASE2018 Task 4 Winner	13
4	Method	14
4.1	Introduction to the Method	14
4.2	Required Data	16
4.3	Data Collection	16
4.3.1	Data collected during project vs ExtraSensory dataset	18
4.4	Data Pre-Processing	19
4.4.1	Audio	19
4.4.2	GPS	20
4.4.3	IMU	20
4.4.4	Audio	21
4.5	Data Training, Testing and Evaluation Set	21
4.5.1	“kfold” Cross validation	21
4.6	Classifier	22
4.6.1	Model Selection	24
4.7	Model Validation	26
4.8	Terminology	26
4.8.1	Hyperparameter Optimisation	28
4.9	Machine Learning Libraries Used	32
4.10	Method Summary	33
5	Results	34
5.1	Binary MLPClassifier (Audio Only)	35
5.2	Binary MLPClassifier (Audio, Accelerometer, Gyroscope, GPS)	35
6	Discussion	35

7 Conclusion	37
8 References and Citations	38
A Full Extracts of Journals	40
B Reflection	40
B.1 Lessons Learnt	40
C Dataset Feature Importance	41
D Trade-Review of "Online vs Offline" processing	41
E GridSearch Attempts	41
F Table of ExtraSensory Labels	42
G Writing you're not sure if you want to keep	42
G.1 Funny Future Work	42
G.2 Development Tools and Technologies	42
G.3 Discussion on practical aspects of IMU and Audio	43
G.4 Nearest Neighbour	43

1 Introduction

This project aims to answer the question “[What is] the effect of supplementary data on Acoustic Event Classification through Machine Learning”. To do this, results will analyse the F1 Scores of a binary MLPClassifier, and a multi-label MLPClassifier, both trained on the ExtraSensory dataset using the exact method described in Section 4. Additionally, the output of the classifiers will be analysed for “validity”. Does the classifier successfully detect and classify an Acoustic Event?

The desired results are a demonstratable difference (or lack thereof) in the F1 score, after comparing a classifier with and without supplementary data. This will

Debriefing is a tool used to educate an individual, inform a collective and create a holistic historical recording of an event. Wikipedia defines the process of debriefing as: (1) Receiving an explanation, (2) Receiving information and situation-based reminders of context, (3) Reporting of measures of performance, and/or opportunities to further investigate the results of a study, investigation, or assessment of performance after participation in an immersive activity is complete [Wik \(????\)](#)^[FLAGGED FOR REVIEW]. Traditionally this is done through a collection of primary sources (where available) and secondary sources. In the information age, these primary sources have expanded to include large datasets, recordings and other digital forensics. This new and increasing range of primary sources provides the potential to significantly improve the steps of debriefing, and ultimately improve debriefing outcomes.

This paper will report on the Author’s progress of harnessing this new potential and is organised as follows: section 2 will discuss how the scope has been analysed and reduced, section 3 will review the literature of this research topic, section 4 will discuss the method of implementing the technology, section 5 will discuss the evaluation methodology, and section 6 will discuss the results of the proof-of-concept.

2 Problem Analysis

This project approached the problem through a System Engineering perspective, a process which “is a structured and systematic methodology providing greater visibility and control over . . . new system[s]” [Limited \(2018\)](#). This perspective was informed by ISO/IEC/IEEE 15288 [for Standardization \(2015\)](#) and Burge Hughes Walsh’s System Engineering Toolbox [Limited \(2015\)](#). The tools adapted and applied throughout this section are primarily derived from that “System Engineering Toolbox”.

The purpose of Problem Analysis is to define the scope of this project to be achievable, measurable and practical to implement. This took 2-4 weeks through June 2018, and remains an ongoing task.

The following System Engineering tools have been adapted and implemented to develop this paper and the author’s understanding of the problem.

2.1 Current 18 Words

A tool called “18 Words” was used to constantly refine and maintain a description on the scope as it changed throughout problem analysis. The current “18 Words” is the following: “[The project is] a portable proprietary format compliant track file data recording system coupled with offline processing to log specific physical events in proximity that occurred during recording.”

2.2 Tree Diagram

With the new understanding of what the project is, a Tree Diagram was drafted to explore missed requirements, hidden modules and other aspects of the project not yet considered. Figure 1 shows a Tree Diagram breakdown for this FYP.

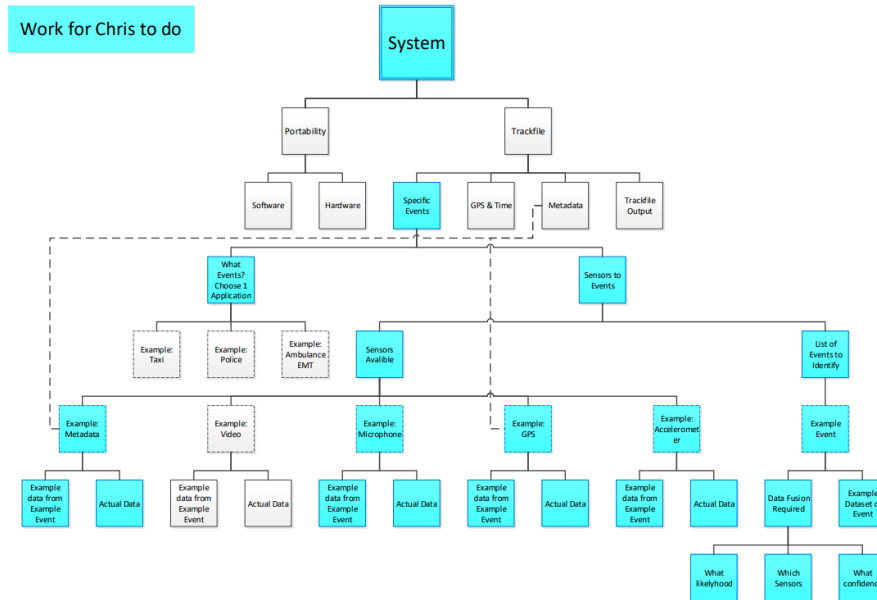


Figure 1: A tree diagram for the project, being used as a method of allocating work

2.3 Quad of Aims

The Quad of Aims is a tool used to explore 4 critical, high level aspects of the project as explained in table 1, and shown in Figure 2.

Table 1: Quad of Aims

Label	Description
Purpose	This is our “18 words”
Stakeholders	University, Author’s Work
Deliverables	Documentation, recommendations, FYP report, proof-of-concept
Measure of Success	Review of Documentation by SMEs, review of FYP report, dry and wet run of proof-of-concept

This is typically done to instigate early works on a project.

2.4 Input Output Analysis

The Input Output Analysis of the system informs the bounds and requirements to operate the system. In this situation, it helped consider the full scope of the project; this includes the technical and non-technical aspects of undertaking the MECH4841 Project as shown in Figure 3.



Figure 2: A Quad of Aims for the project might be A3 in size and would have the relevant information embedded. It may also be completed on a white-board.

2.5 Affinity Diagram

An Affinity Diagram helps to solidify the high level, vague ideas, requirements and tests into a more detailed view. There is a focus on putting tangible measurements onto requirements. During this process, system architecture decisions were made such as a trade review into online vs offline processing, and subsequently splitting the system into smaller modules. Our primary system had 2 key sub-systems: The recording, and the processing systems. An Affinity Diagram of the top-level system was created, and the 2 affinity diagrams of the sub-systems were informed and made from this. After this was complete, previous work was updated to reflect these changes. Figure 4. [\[citation needed\]](#)

2.6 Systems Map

A Systems Map takes the Affinity Diagram, Input Output Analysis, and Tree-Diagram to identify the processes inside sub-modules that are needed to design the system. This was used to great effect to measure and estimate the workload necessary to implement data fusion alongside deep learning of acoustic classification and detection as shown in Figure 5.

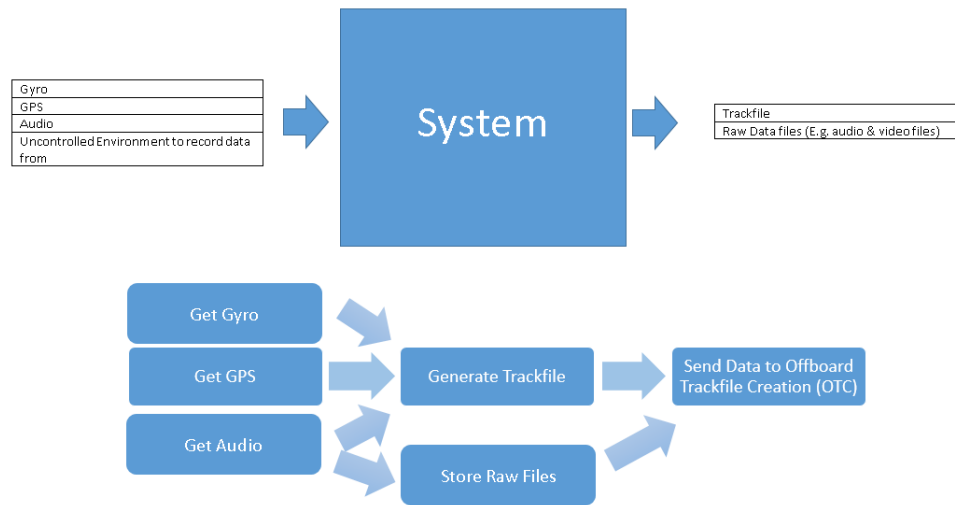


Figure 3: The input output analysis for the system

2.7 Sequence Diagram

A Sequence Diagram was developed to analyse the flow of data through the processes identified in the System Map. This is shown in Figure 6. This helped to predict and manage any potential complexities and logistics related to the specific needs of each process. Originally, This Sequence Diagram was used to justify removing data fusion from the scope due to the large workload required to implement alongside machine learning. Later in the project, it was brought back into the scope of the project to help compare and evaluate it's effectiveness in improving the overall system.

2.8 N² Analysis

An N² Analysis methodically expands on the what data moves around the system. This is to complement the discussed complexities in the Sequence Diagram by documenting what data is expected. A example of this is shown in Figure 7

2.9 Spray Diagram

The Spray Diagram shown in Figure 8 shows how details of the system can have multi-factored effects in design requirements, outcomes and operational use. Of interest in the diagram is the relationship between high-quality output will require high-quality input, and this may increase production and design costs.

2.10 Matrix Diagram

A matrix diagram was made to review the now more refined, reduced scope as shown in Figure 9. This works by using a “strong”, “weak” or “none” indicator for each aspect of the project. It highlighted the difficulty in balancing the needs of both major stakeholders.

2.11 Final Problem Analysis and Discussion

The purpose of the project is to hypothesis and produce a proof-of-concept of how digital technologies could improve debriefing processes and outcomes. This seems achievable within 1 year, and should

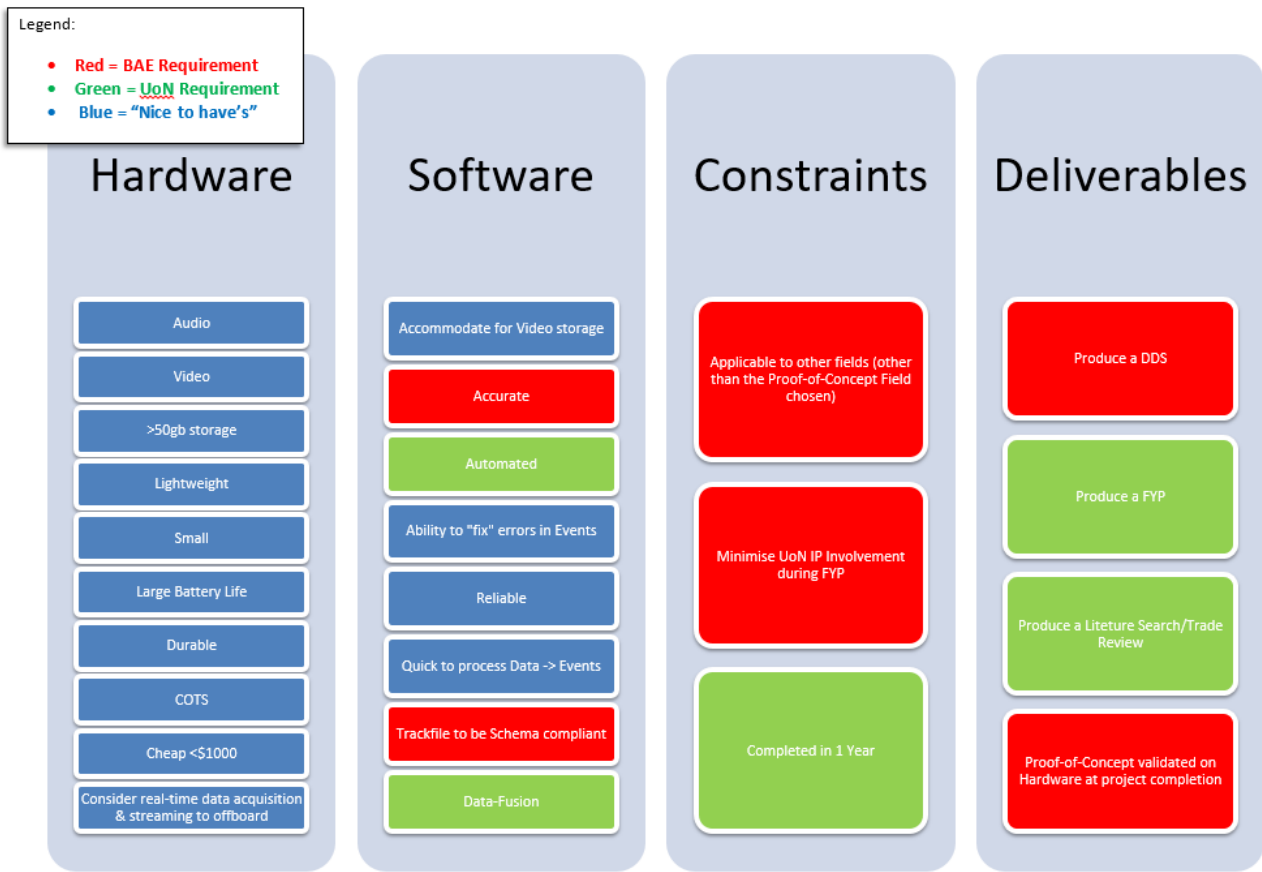


Figure 4: The Affinity Diagrams for the full system architecture

provide an excellent learning opportunity for the author.

3 Review Of Literature

A literature review of the current research was conducted with guidance from University of Queensland's guide to Literature reviews [of Queensland \(2018\)](#) . There are 3 nested areas of research that are of interest to this project. The wider classification and regression research topic, the machine learning topic that covers many various fields of implemented and theoretical machine learning applications, and finally the specific field of Acoustic Event Detection and Classification (AED/AEC); this field has traditionally relied on older techniques, and is undergoing a revolution with machine learning.

3.1 Machine Learning

A definition for Machine learning is taken from Tom M. Mitchell's 1997 textbook, "Machine Learning"

A computer program is said to learn from experience E with respect to some task T and some performance measure P , if its performance on T , as measured by P , improves with experience E . [Mitchell \(1997\)](#)

This definition of Machine Learning can be applied to this project; the task T (AED/C), the E (Audio,

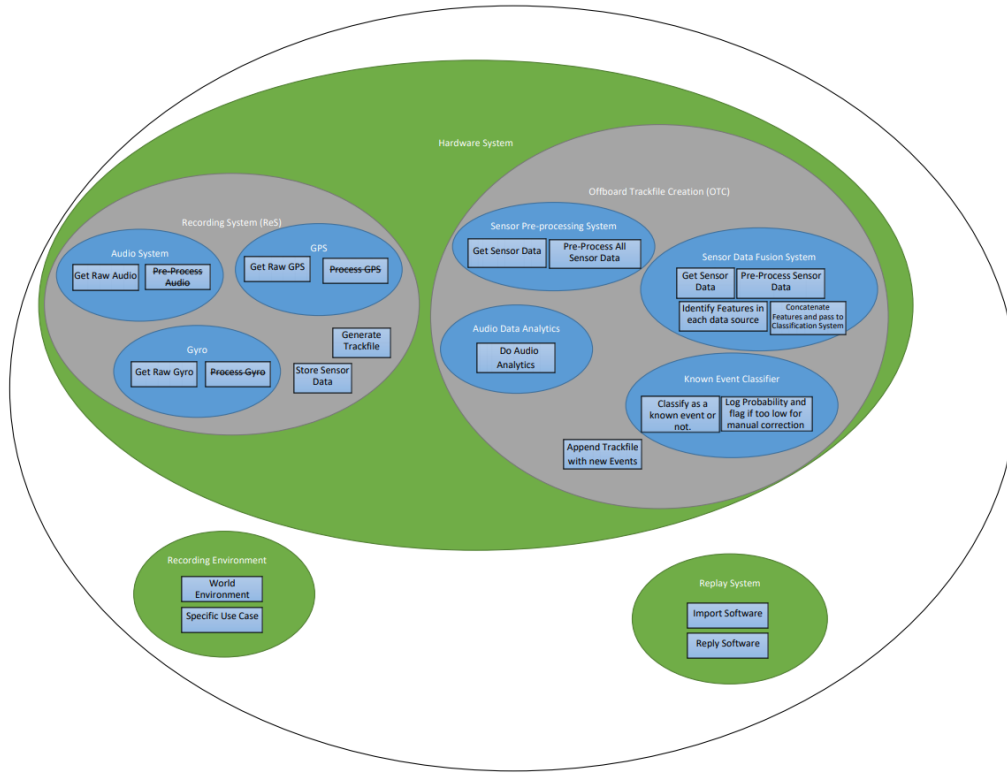


Figure 5: A systems map for the project

IMU, GPS), and P (a traditional F1 score).

The computer program described in this project has been produced in Python, using the sklearn and TensorFlow libraries. [FLAGGED FOR REVIEW]

3.2 Acoustic Event Detection and Acoustic Event Classification

Historically the technologies for AED/AEC have been Support vector machine (SVMs), Hidden Markov Models (HHMs), and more generally, basic DSP (Digital Signal Processing) classifiers. But in the last 8 years, AED and AEC research has emerged as a hot topic of research. Research focuses on the two key tasks: Detection (when did an acoustic event occur in the audio, and when did it stop?) and classification (what sound occurred?). New research papers are often the result of DCASE competitions, which is an official IEEE Audio and Acoustic Signal Processing (AASP) competition. Figure 11.

3.3 DCASE Competitions (Detection and Classification of Acoustic Scenes and Events)

The DCASE competitions format started 2013, had its second competition in 2016, and since then has been a yearly event. Since 2017, there has been a new format where the competition starts march, formal results are released by September, and a workshop for participants on the best team's work in November [DCASE](#). (???). 2016 was the first to feature Machine Learning (winning teams incorporated machine learning into ensemble models/classifiers). By 2017 all entries utilised machine learning. Recently the results from DCASE2018 have been published; progress in the AED/AEC field has allowed more sophisticated, real-world applications to be evaluated (Challenge Task #4) [201](#)

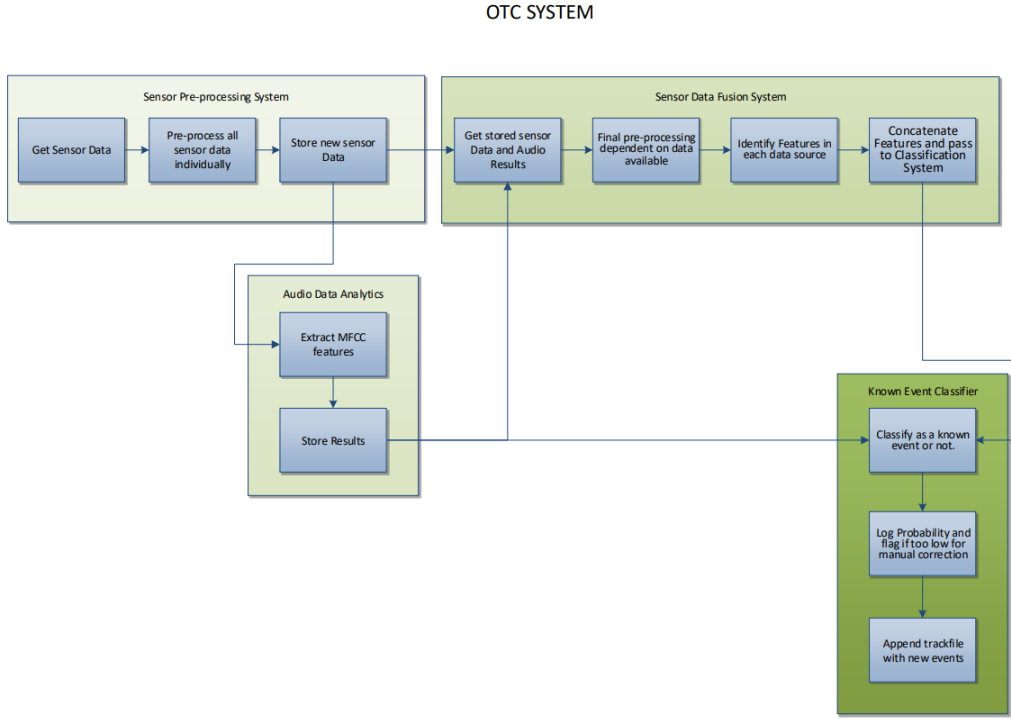


Figure 6: A basic Sequence Diagram for the project

(2018). The complexity of the dataset, the number of acoustic classes and fidelity of output are unprecedented in the AEC field. The winning results at this level of complexity are not as high quality compared to last year’s more controlled environment/dataset (2017 had a 41.7% F1 score201 (2017) vs 2018’s 32.4%201 (2018)), but pre-sent the best precedent for the problem addressed in this paper.

3.4 Leading Research, DCASE2018 Task 4 Winner

The winning model used in the DCASE2018 Task 4 challenge used a “Mean-teacher” model for classification (useful abstract results averaging applicable to a range of classifier models), a CNN for context gating (a pre-classifier step to improve flaws in training methodologies for some machine learning modelsDauphin et al. (2016)) and a bidirectional recurrent neural network (RNN) to improve the utilisation of unlabelled, unbalanced training datasetsJiaKai (2018); this last component is important for this paper, as the leading dataset outside of DCASE challenges is the weakly labelled Google AudioSet, of which only small percentage of is balanced.

Table 2: AudioSet data on cars

Dataset	Number of videos	Duration hours
Evaluation	280	0.8
Balanced train	296	0.8
Unbalanced train	40,978	113.3
Overall	41,554	114.9

Sensor Pre-processing	Audio pre-processing gets passed onto the Audio Analytics module	Sensors Pre-processing gets passed onto the sensor data fusion module		
	Audio Data Analytics	Processed audio gets passed onto the sensor data fusion system <i>FeedbackLoop</i>	Audio get passed onto the known event classifier	
	Possibility of training data coming back from Known Event Classifier helping train the audio data analytics <i>FeedbackLoop</i>	Sensor Data Fusion System	Fused data gets passed into the known event classifier	
Lessons learnt should inform better sensor pre-processing design	Possibility of training data coming back from Known Event Classifier helping train the audio data analytics		Known Event Classifier	

Figure 7: A N² Analysis for the project

$$CarAudioSet : \frac{0.8}{114.9} = 0.69\%$$

4 Method

4.1 Introduction to the Method

In a traditional machine learning problem there are 5 stages to the method;

Stage 1: Data

Stage 2: Preprocessing

Stage 3: Model

Stage 4: Fitting/Training

Stage 5: Evaluation

[citation needed]

Stage 1 Stage 1 is choosing what raw data is available and what information the model must product. This is the inputs and the outputs that you want from the system, and the expected result.

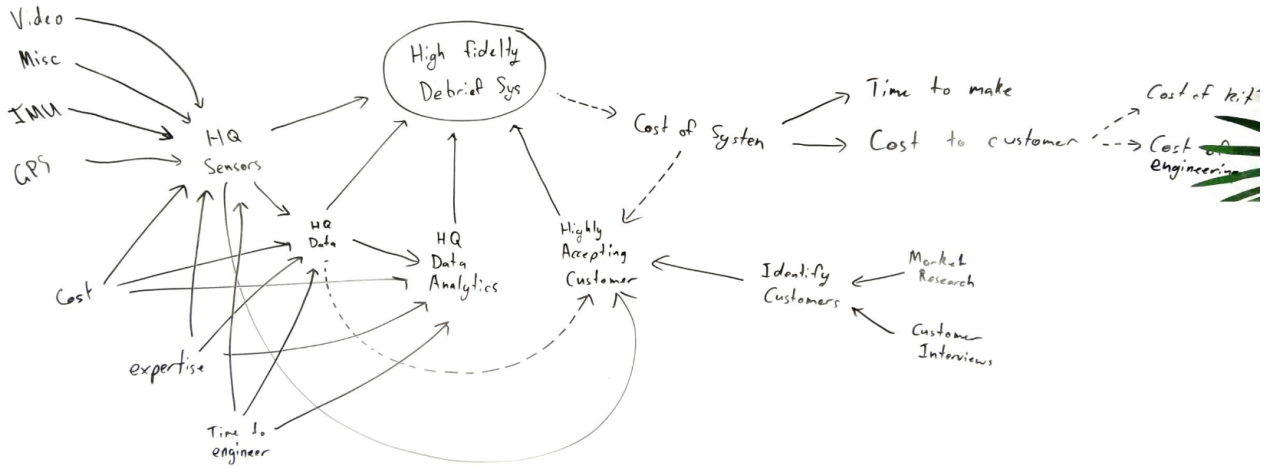


Figure 8: A Spray Diagram for the project

Stage 2 Stage 2 is preprocessing. This is about taking the data available transforming it. This could be optimising, filtering, feature building (producing refined data of interest to the model e.g. MFCCs from raw Audio) and in many other options. Preliminary features are also evaluated in this stage to ascertain their quality and effectiveness in a preliminary model.

Stage 3 Stage 3 is picking the model. The method of picking a model may be based on doing research into what is best practice for your particular problem, looking at recent papers and investigating whether the particular topic demands any special requirements. Once a model is picked, the parameters of the model must be chosen (or evaluated) before a final model can be selected. This is done in the following two stages.

Stage 4 Stage 4 is fitting the model. The selected algorithm may need to be trained or fitted. This is the process of taking labelled input / output data (denoted as X, y) and supervising the model as it attempts to reproduce y from X . A common method to achieve this is to cost Stochastic Gradient Descent (SGD) and back propagation [LeCun et al. \(2012\)](#). In summary, it's the process of changing the weights proportionally to their contribution to the error. Equation 4.1 adapted from [LeCun et al. \(2012\)](#) demonstrates this, where W is the Machine Learning weights for a given topology, η is the proportional factor or step size, and the partial represents how each weight contributes to error.

$$W(t) = W(t - 1) - \eta \frac{\partial E}{\partial W} \quad (4.1)$$

Once trained on a training set, the next stage is validating whether the model is sufficiently fit.

Stage 5 Stage 5 is validation. Validation is the process of checking the performance of the against certain criteria. For instance, a model may be checked for overfitting, underfitting, bias, sensitivity, specificity, accuracy, precision, AUC (area under the (ROC) curve), and F1 score to name a few. F1

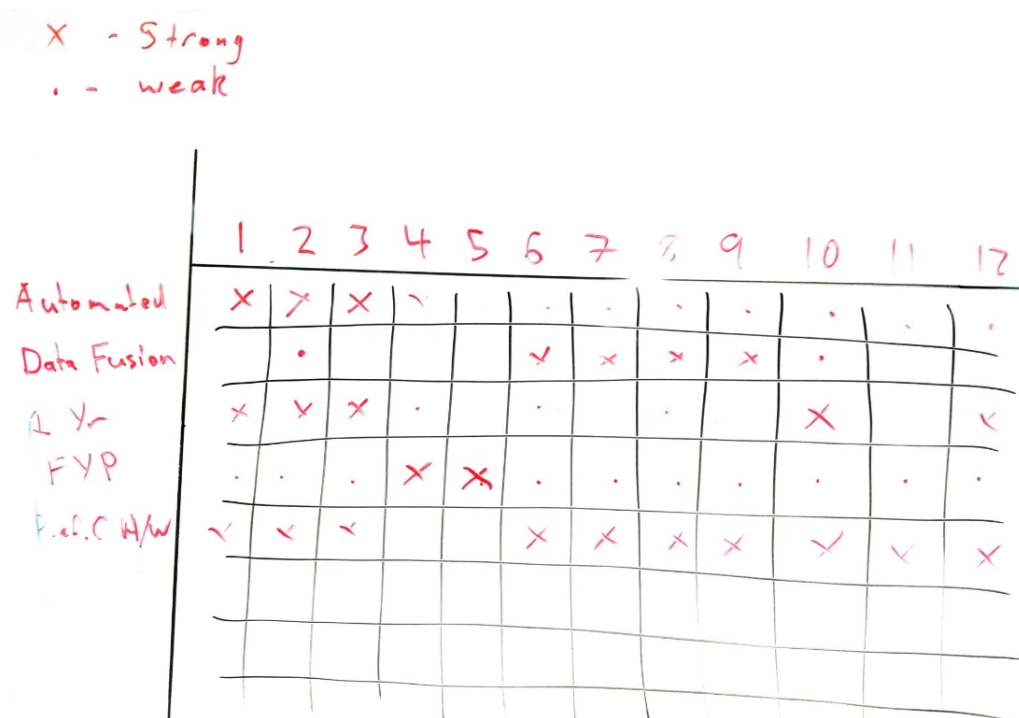


Figure 9: A Matrix Diagram for the project

score will be used as the primary measure because the F1 has been is the research industry standard for evaluating algorithms.

It is beneficial to work up slowly to this goal. external to the method of building the machine learning classifier is the work involved in developing the pipeline associated with it. The main tasks to implement a machine learning classifier are as follows:

- Develop the use-case, and identify the output
- Investigate what data is available, and whether further data/datasets would be beneficial
- Choose a framework/technology to implement the machine learning algorithm in
- Build the machine learning classifier as described above, and test it's effectiveness
- Validate the whole system by using the classifier in the pipeline (by take the classified output, and appending it to the trackfile)

4.2 Required Data

The Machine Learning uses GPS, accelerometer, gyroscope, and audio. A Google Pixel 2 was selected as the model data acquisition, which table /ref below shows.

4.3 Data Collection

The choice of what data to collected and use was found through the first problem analysis step. As the goal is to maximise the accuracy of the classifier (whilst still being an achievable goal), the data

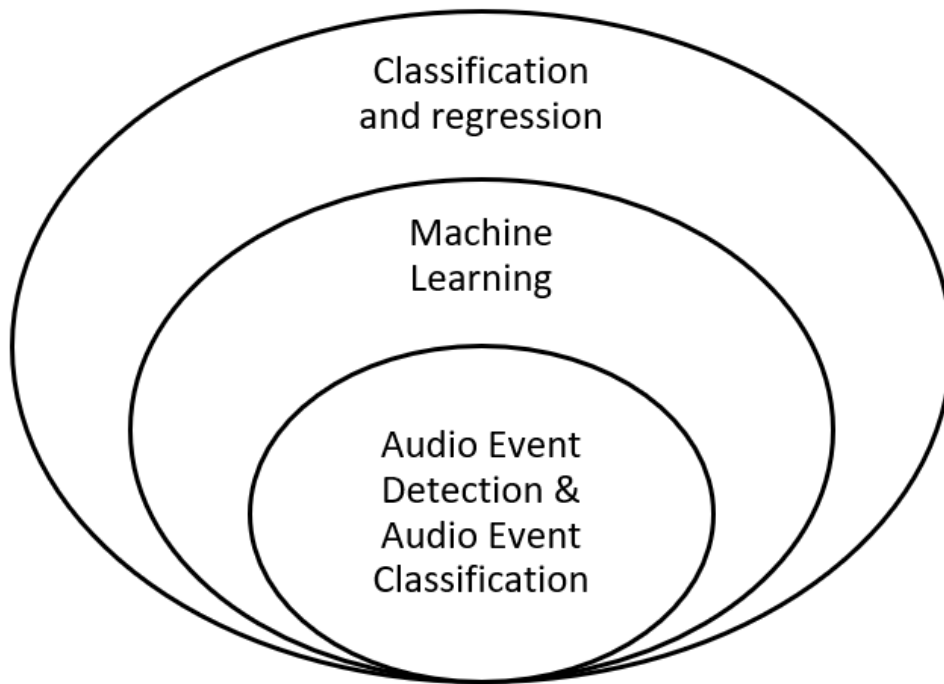


Figure 10: The literature for this project is a niche topic in the Machine Learning area of research

Table 3: Data Available [Havard \(2016\)](#)

Sensor	Description
GPS	The Pixel 2 uses the Snapdragon 835 System on Chip GPS receiver
Accelerometer / Gyroscope	The Pixel 2 uses a Bosch BMI160
Audio	The Pixel 2 uses a Qualcomm WCD9395 Audio Codec chip and a NXP TFA9891UK Audio Amplifier

chosen for use and collection was GPS, accelerometer, gyroscope, and audio. As described in the problem analysis phase, the best platform for this would be a smartphone as shown in section 2 and appendix D.

Android was chosen as the platform as it had significant support for reading from the on-board sensors. The data available from the Android API can be seen in table 4, of which GPS and acceleration will be used. It was also chosen due to it's support for C/C++ through the Native Development Kit (NDK). The programming language C is required knowledge for many courses in the University of Newcastle Bachelor of Mechatronics degree.

The app development took 5 weeks to get to a useful state. It is able to record Audio to a .wav file and record raw GPS lat/long/alt, raw accelerometer, and raw gyroscope to a csv. There are technical issues still present in the app (notably that the screen must be unlocked to record).

During this time, an Atlassian JIRA environment was set up to facilitate management (using concepts from ENG3500 - Project Management). The Atlassian "Git-Flow" methodology was also learnt and used.

I helped with 30% of the overall programming, and code reviewed the product along with other colleagues. Credit goes to [Sebastian Wallman].




Figure 11: DCASE Researcher Toni Heittola’s F1 Results reflect the AED/AEC field of research with drastic improvements over the last 8 years.[Heittola \(2018\)](#)

4.3.1 Data collected during project vs ExtraSensory dataset

The table ?? shows a snippet of the Comma-Separated Values (csv) data produced from the app. The id is a unique count, The table ?? shows a snippet of the csv data from the ExtraSensory dataset.

Both datasets share the same timestamp format, so a “Time” column has been added to the table for illustration purposes only.

The ExtraSensory dataset is compiled from 60 users with a diverse range of phones (“34 iPhone users, 26 Android users.”[Vaizman et al. \(2017\)](#)). This is in contrast to this project’s phone range, which was limited to 2 Android devices, a Samsung S6 and a Lenovo Zuk Z2. The Samsung uses a InvenSense MPU-6500, the Lenovo Zuk Z2 unknown.

id	attr_time	Time	attr_x	attr_y	attr_z
1	1547182324433	2018-01-11 14:52	-0.229843	3.1795	8.52336
2	1547182324442	2018-01-11 14:52	-0.257614	2.93185	8.61444
3	1547182324451	2018-01-11 14:52	-0.179626	2.81464	8.81348
4	1547182324467	2018-01-11 14:52	-0.118927	2.82674	8.70142
5	1547182324498	2018-01-11 14:52	-0.191544	2.95287	8.66815

Listing 1: Code snippet for getting Accelerometer data from the Android System as above [Google \(2019\)](#) above

```

1 public void onSensorChanged(SensorEvent event) {
2     // Get X,Y,Z values
3     accX = event.values[0];
4     accY = event.values[1];
5     accZ = event.values[2];
6
7     ...
8 }

```

timestamp	Time	raw_acc:mean	raw_acc:std	raw_acc:moment3	raw_acc:moment4
1464129912	2016-05-24 22:45	1.011438	0.012573	0.023013	0.04124
1464129950	2016-05-24 22:45	1.011233	0.009356	-0.005622	0.016687
1464130031	2016-05-24 22:47	1.013422	0.018068	-0.008593	0.039286
1464130109	2016-05-24 22:48	1.014891	0.0164	0.021383	0.038825
1464130130	2016-05-24 22:48	1.017487	0.022632	-0.012891	0.037226

“Devices: The users in ExtraSensory had a variety of phone devices. iPhone generations: 4, 4S, 5, 5S, 5C, 6 and 6S. iPhone operating system versions ranging from iOS-7 to iOS-9. Android devices: Samsung, Nexus, HTC, moto G, LG, Motorola, One Plus One, Sony. ”

4.4 Data Pre-Processing

Python was used to process the data before the Machine Learning modules of code, and Matlab was used with GPS, IMU to prototype the algorithms.

The data from the ExtraSensory dataset is provided in it's raw format, which is best option is for the data to be raw, as this allows platform independent filtering, smoothing and sensor fusion [\[citation needed\]](#).

4.4.1 Audio

Data preprocessing of the Audio was done by Mel Frequency Cepstral Coefficients (MFCCs). Start with a recorded .wav sample, and split the whole sample into smaller sections via the sliding window method [\[citation needed\]](#). Take each window and calculate the Fourier transform. This in turn is both “Feature Engineering”, as well as preprocessing to be training compatible.

$$WindowSample(t) \xrightarrow{\mathcal{F}} WindowSample(\omega) \quad (4.2)$$

Audio to a .wav file and record raw GPS lat/long/alt, raw accelerometer, and raw gyroscope to a csv.

“First, resample the audio clips at 22,050 Hz, because the high frequency part of sound signal is not useful for event detection in daily life

Extract the log mel-spectrogram from the audio clips by 128-bin, 2048-window and 365-hop (1683-overlap)” Aka classic sliding window, but this specifies some good over-lap, data density etc.

Listing 2: Matlab trial of MFCC extraction

```

1 speech=readwav(file_path,'s',-1);
2 %rng('default');
3 %speech=speech+randn(size(speech))*eps; %dithering
4 %----- PRE-EMPHASIS -----
5 speech = filter([1 -0.97], 1, speech);
6 %----- FRAMING & WINDOWING -----
7 frame_length_inSample=(Fs/1000)*Window_Length;
8 framedspeech=buffer(speech,frame_length_inSample,frame_length_inSample/2,'nodelay');
9 w=hamming(frame_length_inSample);
10 y_framed=framedspeech.*repmat(w',size(framedspeech,1),1);
11 %-----
12 f=(Fs/2)*linspace(0,1,NFFT/2+1);
13 fmel=2595*log10(1+f./700); % CONVERTING TO MEL SCALE
14 fmelmax=max(fmel);
15 fmelmin=min(fmel);
16 filbandwidthsmel=linspace(fmelmin,fmelmax,No_Filter+2);
17 filbandwidthsf=700*(10.^(filbandwidthsmel/2595)-1);
18 fr_all=(abs(fft(y_framed',NFFT))).^2;
19 fa_all=fr_all(1:(NFFT/2)+1,:);
20 filterbank=zeros((NFFT/2)+1,No_Filter);
21 for i=1:No_Filter
22     filterbank(:,i)=trimf(f,[filbandwidthsf(i),filbandwidthsf(i+1),...
23         filbandwidthsf(i+2)]);
24 end
25 filbanksum=fa_all*filterbank(1:end,:);
26 %----- Calculate Static Cepstral -----
27 t=dct(log10(filbanksum'+eps));
28 t=(t(1:No_Filter,:));
29 stat=t';
30 delta=deltas(stat',3)';
31 double_delta=deltas(delta',3)';
32 %-----

```

4.4.2 GPS

Talk about what a GPS is, and how we get the signal The GPS signal is requested by the "Fused Location Provider", in the Google Play services location API and comes heavily prefiltered^[citation needed]. Further filtering is possible by a Kalman filter with the other physical sensors.

4.4.3 IMU

For the purpose of this project, the data collected through the Android App applied a Low Pass Filter^[citation needed]. A Kalman Filter was considered, but wasn't used to avoid modelling the specific android sensor plant dynamics.

Figure 12




Figure 12: Results of Low Pass Filter vs Raw accelerometer data

4.4.4 Audio

4.5 Data Training, Testing and Evaluation Set

The ExtraSensory dataset contains over 300,000 samples across 60 Users (in a CSV format with 278 feature columns), as shown in ???. Each row contains an a 'snapshot' of the sensors at a frequency of once a minute. There are large non-continuities across each user. The dataset is unbalanced, and has had it's labels cleaned by the researchers involved. Further details about the users themselves can be found in 8, or at it paper [Vaizman et al. \(2017\)](#).

The dataset was considered for use in 2 ways; training on a per user instance, or training on a per sample instance. The benefits of per user is to isolate recording hardware, and other plant noise. However this is undesirable, as this can result in a low variance in ϵ (bias), which increases the likelihood of the learnt function (\hat{f}) to overfit the signal of the training users ($\hat{f} = f(x) + \epsilon$). This, to minimise it's impact in evaluation. Thus is the benefit of training on a per sample basis. By introducing other user noise into a training set, we minimise overfitting "for free". Finally as a "kfold" (or other Cross-Validate methods) is applied to the whole training dataset for training + validation purposes, the likelihood of the random selection incorporating a larger variance of noise increases. As a result, the data is used on as a per sample basis.

4.5.1 "kfold" Cross validation

kfold Cross Validation can be described as:

"... approach involves randomly dividing the set of observations into k groups, or folds,

Listing 3: Smoothing of Acceleration data

```

1 importURL = "C:\FYP\report\Code\2019-01-11T16-26-00_androidAccData.csv";
2 AccData = csvread(importURL,1,1);
3 SmoothedAccData = zeros(size(AccData,1),4);
4 t = 1:size(AccData,1);
5 lowpassAlpha = 0.40; % 80%
6 for i = 1:size(AccData,1)
7     for j = 2:4
8         % Low Pass Filter on
9         SmoothedAccData(i,j) = SmoothedAccData(i,j) ...
10             + lowpassAlpha * (AccData(i,j) - SmoothedAccData(i,j));
11     end
12 end
13 for plotNum = 1:3
14     subplot(3,1,plotNum);
15     plot(t,AccData(:,plotNum+1),'k');
16     hold on
17     plot(t,SmoothedAccData(:,plotNum+1), 'r', 'LineWidth',1);
18     ylabel('m/s^2');
19     switch(plotNum)
20         case(1)
21             title('Raw_X_Acceleration_vs_LPF_Smoothed');
22             legend('Raw', 'LPF');
23         case(2)
24             title('Raw_Y_Acceleration_vs_LPF_Smoothed');
25         case(3)
26             title('Raw_Z_Acceleration_vs_LPF_Smoothed');
27             xlabel('Sample_Number');
28     end
29 end
30 set(gcf,'color','white')
31 saveas(gcf,'AndroidDataExample.png')

```

of approximately equal size. The first fold is treated as a validation set, and the method is fit on the remaining $k - 1$ folds. ”

[James et al. \(2017\)](#) A key element expanding on this is the ability to cycle through ‘k folds’, changing which group is the validation, or test set, and training on the rest.

The question of whether supplementary data will improve the classifier’s F1 score is best answered by deciding on an appropriate classifier, and testing it with and without the supplementary data. As discussed, the appropriate classifier could be chosen via hyperparameter testing against a validation set, however the feature analysis carried out above will also influence classifier design. The unbalanced nature of the dataset is also a factor.

4.6 Classifier

A classifier by definition classifies things. A more thorough understanding is that a classifier takes some input data and maps it to an output category, or label. In a Machine Learning context, this mapping is done via the \hat{f} learnt function. For this project the classifier is a learnt function that takes

Table 4: Data Available^[citation needed]

Sensor	Description
TYPE ACCELEROMETER	Acceleration force along the x axis (including gravity). Acceleration force along the y axis (including gravity). Acceleration force along the z axis (including gravity).
TYPE ACCELEROMETER UNCALIBRATED	Measured acceleration along the X axis without any bias compensation. Measured acceleration along the Y axis without any bias compensation. Measured acceleration along the Z axis without any bias compensation. Measured acceleration along the X axis with estimated bias compensation. Measured acceleration along the Y axis with estimated bias compensation. Measured acceleration along the Z axis with estimated bias compensation.
TYPE GRAVITY	Force of gravity along the x axis. Force of gravity along the y axis. Force of gravity along the z axis.
TYPE GYROSCOPE	Rate of rotation around the x axis. Rate of rotation around the y axis. Rate of rotation around the z axis.
TYPE GYROSCOPE UNCALIBRATED	Rate of rotation (without drift compensation) around the x axis. Rate of rotation (without drift compensation) around the y axis. Rate of rotation (without drift compensation) around the z axis. Estimated drift around the x axis. Estimated drift around the y axis. Estimated drift around the z axis.
TYPE LINEAR ACCELERATION	Acceleration force along the x axis (excluding gravity). Acceleration force along the y axis (excluding gravity). Acceleration force along the z axis (excluding gravity).
TYPE ROTATION VECTOR	Rotation vector component along the x axis ($x \sin(o/2)$). Rotation vector component along the y axis ($y \sin(o/2)$). Rotation vector component along the z axis ($z \sin(o/2)$). Scalar component of the rotation vector ($(\cos(o/2)).1$
TYPE SIGNIFICANT MOTION	N/A
TYPE STEP COUNTER	Number of steps taken by the user since the last reboot while the sensor was acti
TYPE STEP DETECTOR	N/A

audio, accelerometer, gyroscope and GPS data and maps that data to a label one of the labels found in Appendix F, table 9.

Binary, Multi-class, and Multi-Label classifiers If a classifier is to map data to a label, it stands to reason that the label must be well defined. A binary classifier as the name suggests will categorise input as 1, or 0 against a label. This could be result in "Car" or "Not Car", but it can also be extended to multiple classes, for example "Not Car", "Holden", "Ford" etc. This can be done via a ensemble of binary classifiers (combing together the output of the same dataset that has been parsed into a "Ford" classifier, a "Holden" classifier etc.), or more recently through a "One Hot"^{Huffman (1954)} encoded output array as seen in table 6. A key element of multi-class classifiers are that the classes are mutually exclusive. The car brands in 6 are mutually exclusive.

However there is another type of classifier called the multi-label classifier. An example of a multi label classifier is this project. The project may want to label whether signal in the provided data represents jogging, walking, talking, eating etc. For these labels, they are not mutually exclusive and could all be occurring at the same time.

Table 5: ExtraSensory Data Breakdown [Vaizman et al. \(2017\)](#)

	Range	Average	(standard deviation)
Age	18-42	24.7	-5.6
Height	145-188	171	-9
Weight	50-93	66	-11
Body	18-32	23	
Labeled	685-9,706	5,139	-2,332
Additional	2-6,218	1,150	-1,246
Average	1.1-9.7	3.8	-1.4
Days	2.9-28.1	7.6	-3.2

Table 6: Example Hot One Encoding

	Holden	Ford	Other
Ford Fiesta	0	1	0
VW Golf	0	0	1
Holden Captiva	1	0	0

For this project, a binary classifier and a multi-label classifier will be used and evaluated. The reason for this is because the intent of the project is to answer the question does supplementary data assist in classifying an Acoustic Event Detection and Classification (AEC/D) problem, and a binary classifier is a valid and succinct way to answer this question. However for most realistic use cases, what is important is whether the result is valid in a multi-label classifier. For this reason, the methodology section will be applied to creating a multi-label variant of the network. This will use the same hyperparameter optimisation as the binary classifier and will be discussed in the results.

4.6.1 Model Selection

Model selection is an important aspect of Machine Learning, because of what is known as a the "No Free Lunch" theorem.

"No Free Lunch theorems have shown that learning algorithms cannot be universally good. [Magdon-Ismail \(2000\)](#)"

This theorem shows that selection of a Model is non-trivial, and the use case must be factored in.

The model required for the project will have to take into consideration several factors; the type, format and datapoints of data in use, the type and format of the desired information, as well as some unique factors: desire to run onboard, desire to run generalise well (even into different use cases). The standard approach is to either a) try many options or b) have an educated guess as to the

Convolution Neural Networks CNN convolution neural networks are primarily designed and received the best results for image-based classification problems this is because the architecture has been designed with computer vision in mind for for our project

Recursive Neural Networks The recursive neural network (RNN) is not well suited to tabular data this is because and this will not be used because it has been designed for producing multi-class

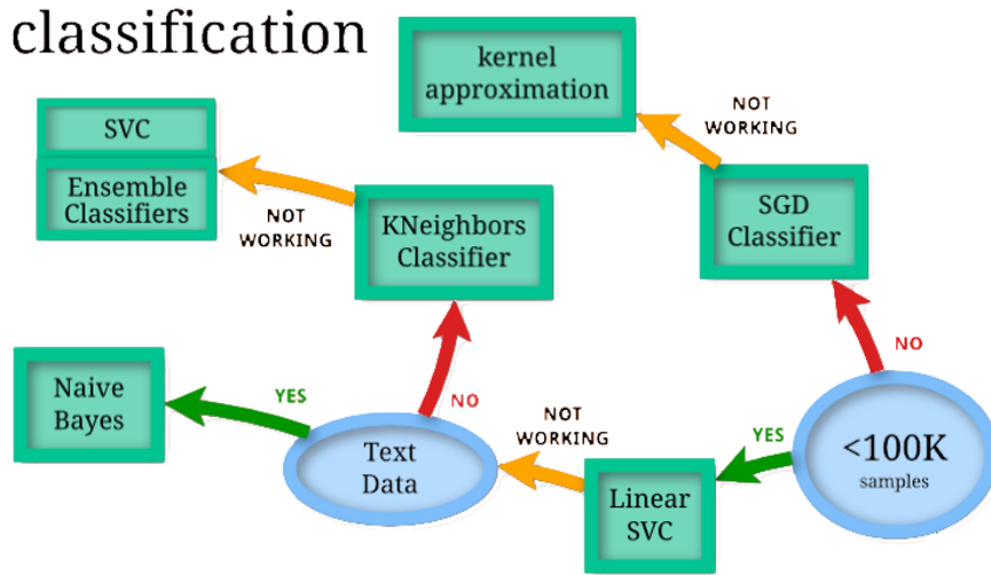


Figure 13: Scikit-learn's suggested model selection process

probabilities we are we could use a recursive neural network to generate a probability output that is to say use it as a regression out progressive algorithm to output all a probabilities however we are after a binary classification and fast it is not relevant more important what the classifier deems as an appropriate probability before it classifies for future work this may in fact be a very useful extension as the tailoring of the sensitivity and specificity probability thresholds of the recursive neural network. This could also be used in conjunction with a soft Max layer^[FLAGGED FOR REVIEW] to produce a normalised a labelling classifier this However it would not be used due to it's extra complexity which does not which is not needed by a use case other options include a random Forest classifier a random Forest classifier is not well suited to unbalanced data sets and would prevent a much bigger problem with for training the future the random forest classifier has been evaluated for comparison in the results section however as a architectural decision it will not be used simile so scikit-learn also has the algorithm extra trees classifier citation needed which can be used in a similar way to the random Forest Hill that you eat better handles label caissons which is important for our use case ll use case is evaluating a binary class against the potential list

Mulilayer Perceptron Feed Forward Neural Network The multilayer perceptron classifier has been selected this

For our project we are using tabular time-invariant data using a snapshot data in a tabular format means that we will not see the benefits of a convolution neural network in a deep learning environment vs the more simple mlpclassifier which trains faster will receive the better result and is expected to compress better winNetwork(parametee will also

Final Model ^[FLAGGED FOR REVIEW]fix the details The MLPClassifier is a feedforward neural network, trained via backpropegation. It was set up using the "lbfgs" solver, an alpha equal to $1e - 5$, and was handselected to have 80 by 10 hidden layers (80 long, 10 deep). This was the classifier to use for evaluation.

4.7 Model Validation

Validation is the process of optimisation the hyperparameters of a model on the hold-out validation set of data. This is an important aspect of machine learning process because model validation cannot be learnt from data. That is to say, hyperparameter optimisation is out of the scope of training via traditional gradient decent and back propagation. Due to the "No Free Lunch" theorem, a model will be chosen and must be optimised before training begins. This is an extension of trying to find the optimal model for the solution.

Model Performance Metrics Before the hyperparameters can be optimised, a metrics to optimise the network on must be selected, of which there are many ways to do it which will now discussed.

4.8 Terminology

Before evaluating metrics, the Terminology of the following formula must be formalised. The first is what a "correct" classification is. The following terminology will be used:

- True Positives (TP) - A correctly classified positive sample
- True Negatives (TN) - A correctly classified negative sample
- False Positives (FP) - A negative example incorrectly classified as positive
- False Negatives (FN) - A positive example incorrectly classified as negative

These four terms are the main parameters in the following formula

Accuracy Accuracy is the measure of how many samples the classifier correctly classified across it's operation, as shown in 4.3.

$$Accuracy = \frac{TP + TN}{TP + FP + TN + FN} \quad (4.3)$$

Accuracy is used when a dataset is balanced (or close to balanced), and when there's no specific weighting on False Positives or False Negatives. Situations where you may have specific weightings on False Positive and False Negatives rates could be the medical industry, where a false positive will have a cost associated with it (the cost of further tests, or screenings) vs a false negative (where by a patient's health outcomes could be severely impacted by the incorrect diagnosis).

In this project, accuracy would be a very poor metric for machine learning because the ExtraSensory dataset compromises 51 labels. When performing a binary classification on any individual label the majority class will be blank/negative. As such, a classifier which would classify every sample as blank/negative would still achieve a high accuracy rate (due to achieving a high True Negative rate).

Confusion Matrix A confusion matrix is a 4 cell table that lists TP, FP, FN, TN. This is a great visual tool to aid human evaluation of the model, as shown in Figure 14. The matrix can be normalised across the rows, as to show what portion of samples were correctly classified. The capacity to directly penalising (and optimise about) the False Positives, and False Negatives, could be an essential tool for a project.




Figure 14: A confusion matrix of the Iris Database. [Pedregosa et al. \(2019a\)](#)

In this project, the Confusion Matrix gets generated to sanity check, and contextualise, other performance measures. It is an important tool for weighing up the validity of classifier before accepting it for real world use.

Recall (or sensitivity) Building upon the direct use of the 4 factors, TP, FP, FN, TN, one can measure the ability of a classifier to “Recall” the true positive rate of a dataset as defined in 4.4. This is the ratio of predicted positives vs the “True” positive rate.

$$Recall = \frac{TP}{TP + FN} \quad (4.4)$$

In this project, the Recall rate is of importance as the user case is to document an Acoustic event, upon which any False Positives could be disregarded by the user, or provide interesting new insight into the behaviour observed.

Precision In contrast to Recall, precision is the measure of the error in Positives classification. That is, out of the samples classified as Positive, how many were actually positive, or the Ratio of correct predictions vs all predictions.

$$Precision = \frac{TP}{TP + FP} \quad (4.5)$$

In this project, this is less important as discussed above. Precision may be important when a high degree of statistical confidence is needed for a decision, but missed opportunities may not be important.

Specificity Specificity is the compliment of to Recall. It is the ratio of correctly negative classifications vs the total number of negative classifications. This is another important metric for training models that are required to be highly “specific” it the model’s predictions.

$$Precision = \frac{TN}{TN + FP} \quad (4.6)$$

In this project, the specificity will ideally be traded off against the Recall rate. For the reasons already discussed, the risk from false positives are negligible.

F1 Score An F1 Score is an industry standard approach to evaluating a model. It is formally defined as the harmonic mean of the Precision and Recall, as shown in 4.7. This score is a method of evaluating a model’s Recall and Precision with equal weighing, but punishing an imbalance between the two scores.

$$F1Score = \frac{2 \times (Precision \times Recall)}{Precision + Recall} \quad (4.7)$$

This project will use the F1 Score as it’s prime metric for evaluating the model, and the effect of supplementary data on the Acoustic Event Detection/Classification task in question. A further extension of the F1 Score is required to apply to a multi class or multilabel problem. This is often done through a sum of F1 Scores across each class or label. For best precision, a weighted could be applied to sum to best reflect (and optimise) the intent of the classifier.

Area Under the Reciever Operating Characteristics (AUROC) Curve Lastly, the Area Under the Reciever Operating Characteristics (AUROC) is an important measure of the a classifiers overall performance. The AUROC demonstrates how confident a classifier is across all class-classification thresholds.

In a more formal sense, the Reciever Operating Characteristics (ROC) is the relation between the True Positive rate against the False Positive rate, and how it changes when you change the classification decision threshold. To calculate the ROC, a sample of TP and FP is taken for a varying threshold (the threshold needed for the classifier to classify a sample as positive or negative).

In this project, the AUROC will not be extensively used, except for educational purposes. The purpose of a AUROC is to demonstrate, and optimise the confidence in classification based on thresholds applied. The AUROC can be used to demonstrate a quanitative assessment of any particular sample (similar to a Regression problem).

4.8.1 Hyperparameter Optimisation

Machine learning is about learning, and adjusting, the internal parameters of a classifier. A hyperparameter is a parameter that cannot be learnt from backpropegation, as it is external to the classifier. These are parameters such as hidden layer size, training rates, classifier specific strategy, etc. The way to choose these hyperparameters is therefore an optimisation problem. The problem is to maximise the score selected above, by trialing different hyperparameters. The following describe methods of finding the optimal hyperparameters.




Figure 15: A ROC Curve example from Google’s Machine Learning shortcourse [Google Brain \(2019a\)](#)

Manual Search Hyperparameters could be manually selected and adjusted, in a search for a more optimal solution. This may be a valid option for interrogating what each hyperparameter does, and the impact on the classification score. This however would not be a valid solution for finding the optimal hyperparameters.

This was tried and yielded mediocre results from default hyperparameter settings on various algorithms.

Brute Force Search To improve upon the Manual Search, one could attempt to try every option available in your parameter space.

```
for hyperparameter1 = 0 to 100 {
  for hyperparameter1 = 0 to 100 {
    TrainNetwork(hyperparameter1,hyperparameter2)
  }
}
```

This improves on the Manual Search, but is very inefficient ($O(n^2)$), and thus extremely slow.

This was not tried during this project, but highlights the importance of the following algorithms.

Grid Search (Manual Grid Search) If the Brute Force Search method attempts to try every combination, then an improvement on this is to better specify the set from which a combination is formed. This is called a Grid Search. A Grid Search iterates through every combination provided. A 2d example shown in Figure 16 shows a grid is generated based on the hyperparameters intervals selected. A power element of Grid Search is the capacity to select non-equidistant intervals.

Grid Search was used for early implementations across both TensorFlow + Keras, and Scikit-learn.




Figure 16: 3 options for Hyperparameter 1 and 2 have been selected. Each combination will be evaluated, as denoted by the green 'x'.

Random Search (Random Grid Search) A Random Search builds upon the Grid Search, but aims to reduce the computation time required to find an optimal set of hyperparameters. This is especially true in higher dimensions^[citation needed]. The key difference between the two searches, is Random Search will not attempt to try every combination. It will sample random number of combinations, which has been “show[n that] empirically and theoretically, randomly chosen trials are more efficient for hyper-parameter optimization than trials on a grid”[Bergstra and Bengio \(2012\)](#). In part, this is due to a principle known as “Embarrassingly parallel”.

“Embarrassingly parallel”[Maurice Herlihy \(2012\)](#) is a critical concept in Machine Learning that allows the Random and Grid searches to be run in parallel. The concept is, because running a trial of 1 point does not have a dependence on any other point, they can be run in parallel. This is where mass parallel computing (such as the operations of a Graphic Processing Units, or GPU) allows very large hyperparameter optimisation searches to be run in parallel.

Due to the simplicity of the multilayer perceptron model selected, random search was primarily used for hyperparameter optimisation.

Bayesian Hyperparameter Optimisation Bayesian Hyperparameter Optimisation is the principle of allowing a search to be influenced by previous experience. This is applied when the cost of the sample, such as a model’s F1-score from a particular combination of hyperparameters, is computationally expensive to calculate. A more formal description can be found in Brochu’s 2010 paper, “A Tutorial on Bayesian Optimization ...”.

“Bayesian optimization employs the Bayesian technique of setting a prior over the objective function and combining it with evidence to get a posterior function. [Brochu et al. \(2010\)](#)”

Bayesian Optimisation for hyperparameter selection was not tested, as it is suggested it is best suited for optimising more complex networks. Snoek et al. (2012) Brochu et al. (2010) It is however an excellent tie in to the MCH3900 Bayesian lectures, and could constitute a future package of work.

Final Hyperparameter output Below are snippets of the an implemented Random Search for an MLP classifier. Of note is the parameters included are listed in the table 7 below. The hyperparameter options listed are a part of an iterative process where the range was manually narrowed after successive RandomSearches. This is a method to increase precision without high initial ranges that would be inefficient.

Table 7: Random Search Hyperparameters

Hyperparameters	Description	Options				
'learning_rate'	How Δ Weight changes wrt loss gradient	"constant"	"invscaling"	"adaptive"		
'hidden_layer_sizes'	$(H_1, \dots, H_{n-1}, H_n)$ where H is is number of nodes, and n is the number of layers	(145,)	(147,)	(149,)	(151,)	(153,)
'alpha'	L2 Regularisation rate, for penalising feature weights. Helps prevent overfitting	0.09	0.11	0.13	0.15	0.17
'activation'	The model function $f(x)$ to fit to the problem function, where $y = f(\sum(W \times x)) + bias$	"logistic"	"relu"	"Tanh"		

These parameters are then placed into a python Dict, and passed into the scikit RandomizedSearchCV function.

[FLAGGED FOR REVIEW] does it need a title/numbers etc.?

```
parameters={'learning_rate': ["constant", "invscaling", "adaptive"],
            'hidden_layer_sizes': [(145,), (147,), (149,), (151,), (153)],
            'alpha': [0.09, 0.11, 0.13, 0.15, 0.17, 0.19],
            'activation': ["logistic", "relu", "Tanh"]}
mlpc = MLPClassifier(verbose=False, early_stopping=True, learning_rate='adaptive', max_iter=1000)
clf = RandomizedSearchCV(estimator=mlpc, scoring='f1', param_distributions=parameters,
                        n_jobs=-1, verbose=1, cv=3);

clf.fit(X_train, y);
print("F-Score: %.2f" % clf.score(X_train, y))
print(clf.best_params_)
```

and the results of which were

```
F-Score: 0.92
{'solver': 'adam', 'learning_rate': 'invscaling', 'hidden_layer_sizes': (150,),
 'alpha': 0.17, 'activation': 'relu'}
```

Details of a Grid Search can be found in Appendix E.

4.9 Machine Learning Libraries Used

The classification libraries used were Google’s TensorFlow and scikit-learn. Both libraries are python based (supporting both Python 3.0, and Python 2.7) [Pedregosa et al. \(2019b\)](#) [Google Brain \(2019c\)](#)

TensorFlow TensorFlow is an opensource machine learning framework, primarily developed by Google. “It has a comprehensive, flexible ecosystem of tools, libraries and community resources that lets researchers push the state-of-the-art in ML and developers easily build and deploy ML powered applications.”?. TensorFlow in conjunction with Keras was used to prototype the initial pipeline. According to the textbook, “Hands on Machine Learning with Scikit, Keras & TensorFlow”, a pipeline is

“ A sequence of data processing components is called a data pipeline. Pipelines are very common in Machine Learning systems, since there is a lot of data to manipulate and many data transformations to apply [Geron \(2017\)](#). ”

TensorFlow and Keras were selected to initially design the classifier due it’s ability to run on the TensorFlow Lite. This is inline with the motivation of the project to allow widespread use of an AED/C solution. The code snippet 4 shows how a model converts.

Listing 4: Code snippet for converting the classifier model to TensorFlow Lite for mobile operation [Google Brain \(2019b\)](#)

```
1 import tensorflow as tf
2
3 converter = tf.lite.TFLiteConverter.from_saved_model(saved_model_dir)
4 tflite_model = converter.convert()
5 open("converted_model.tflite", "wb").write(tflite_model)
```

The original prototyping for audio only classification with TensorFlow and Keras was to use matplotlib to plot the spectrogram of each audio sample, and use a sequential CNN to process the spectrograms. A code snippet for that design is shown in 5.

Due to the sample size of the dataset (AudioSet’s large database, but smaller balanced datasets for the specific use-cases identified), and the use of GPS, Acceleration and Gyroscope data, the classifier could have been prone to overfitting. [\[citation needed\]](#) Dropout layers are method of regularisation, which is used to help prevent overfitting in machine learning. Dropout layers work by “randomly dropping units (along with their connections) from the neural network during training” ?. The purpose of dropout layers in this classifier is to prevent the classifier from overfitting by reducing the over reliance on any individual neuron, but also to reduce the over reliance on any individual data source. This is important because the Audio feature has the highest fidelity (MFCCs across N-bins), and the classifier could have a heavy bias towards it and as a result overfit. Dropout layers will help minimise this overfit. To test this hypothesis, a Random Forest classifier (estimators = 300) was trained on all users to weigh the features, to help identify the design choices. Appendix C shows each feature importance breakdown, where Figure 17 shows the scikit Random Trees feature importance of the set (Audio, Accelerometer, Gyroscope, GPS).

A Random Tree Classifier was used to evaluate feature importance. Random Tree Classifier’s perform poorly on unbalanced datasets. like on on a the ExtraSensory after random undersampled dataset,

during the design phase of the main classifier. Features were Audio (Blue, 26 bins), Accelerometer (Orange, 26 features), Gyroscope (Red, 26 features), GPS (Green, 17 features)

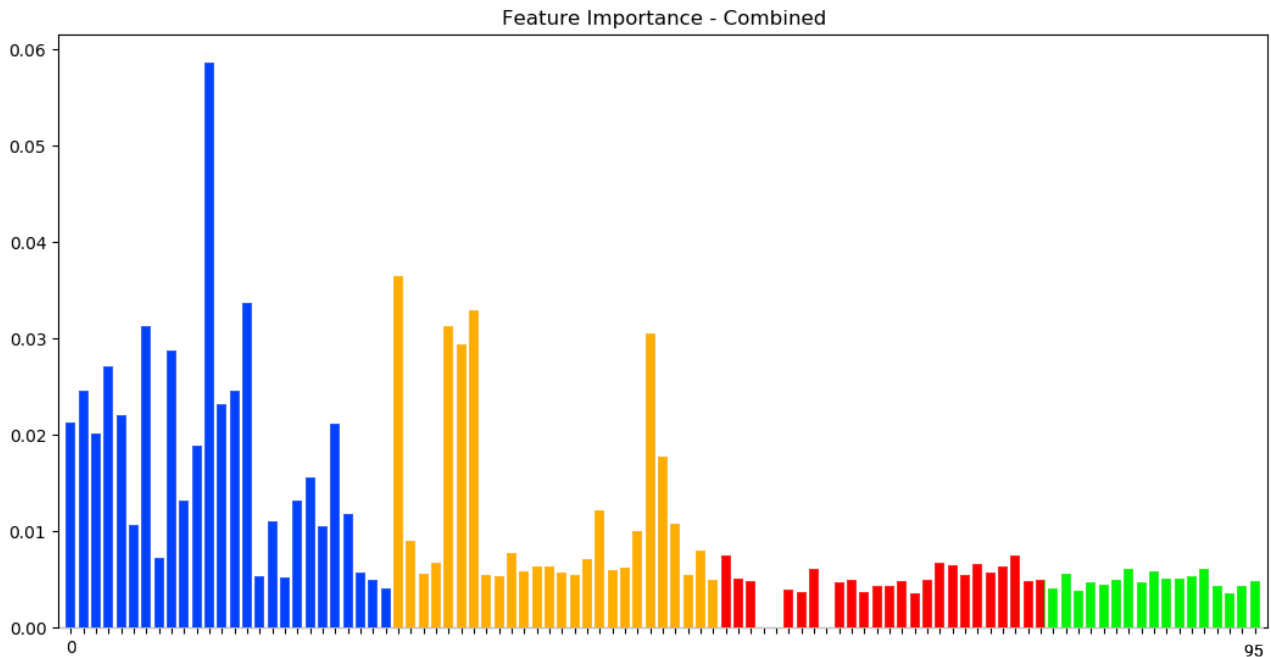


Figure 17: A Random Tree Classifier was used to evaluate feature importance on a the ExtraSensory after random undersampled dataset, during the design phase of the main classifier. Features were Audio (Blue, 26 bins), Accelerometer (Orange, 26 features), Gyroscope (Red, 26 features), GPS (Green, 17 features)

scikit-learn Scikit-learn is an extensive machine learning Python library. The purpose of the library is to remove much of the “boilerplate” code associated with and in implementing machine learning networks and algorithms. It also provides an extensive list of machine learning support functions. When combined with other Python libraries such as numpy, matplotlib and others, it allows for succinct and consistent reproduction of the results this project aims to deduce. It is for this reason that libraries for machine learning have been utilised in contrast with writing a framework from scratch.

4.10 Method Summary

The method details selecting the desired output, the required data input, a metric to score by, a classifier, optimising hyperparameters against the previous options, and finally training and validating against the data.

For this project

1. The desired output is a label (or labels) detailing the Acoustic Event in the sample (e.g. The labels found in [F](#) like ‘Walking’).
2. The required data is the ExtraSensory dataset, or a custom dataset as recorded by the project’s App.

Listing 5: Code snippet for the early CNN model

```

1 model = Sequential()
2 model.add(Conv2D(32, (3, 3), padding='same', input_shape=x_train.shape[1:]))
3 model.add(Activation('relu'))
4 model.add(Conv2D(32, (3, 3)))
5 model.add(Activation('relu'))
6 model.add(MaxPooling2D(pool_size=(2, 2)))
7 model.add(Dropout(0.25))
8
9 model.add(Conv2D(64, (3, 3), padding='same'))
10 model.add(Activation('relu'))
11 model.add(Conv2D(64, (3, 3)))
12 model.add(Activation('relu'))
13 model.add(MaxPooling2D(pool_size=(2, 2)))
14 model.add(Dropout(0.25))
15
16 model.add(Flatten())
17 model.add(Dense(512))
18 model.add(Activation('relu'))
19 model.add(Dropout(0.5))
20 model.add(Dense(num_classes))
21 model.add(Activation('softmax'))

```

3. The score selected is the F1 score. This is due to it being an industry standard.
4. The classifier chosen is the MLPClassifier which is a feedforward neural network, trained via backpropagation.
5. Hyperparameters of the MLPClassifier are chosen via a Random Search Cross Validation.
6. Training and Validation was run with Early Stopping using a 10% validation set, using 59 out of the 60 ExtraSensory users.
7. Final testing is done with the 60th user.

5 Results

The desired results are a demonstratable difference (or lack thereof) in the F1 score, after comparing a classifier with and without supplementary data. This will answer the question, and title of the report, “[What is] the effect of supplementary data on Acoustic Event Classification through Machine Learning”. To this effect, the results will analyse the F1 Scores of a binary MLPClassifier, and a multi-label MLPClassifier, both trained on the ExtraSensory dataset using the exact method described in Section 4. Additionally, the output of the classifiers will be analysed for “validity”. Does the classifier successfully detect and classify an Acoustic Event?

```

--- Training Metrics ---
Samples: 31918 .
Label: 'Walking'
TP: 17641
TN: 14277
Sensors: Aud
Features: 26 features

== Testing with 8023 examples. For label 'Walking' we have 4495 positive and 3528 negative examples.

--- Performance Metrics ---
Accuracy: 0.79
Recall(sensitivity): 0.83
Precision: 0.80
Specificity: 0.73
F-Score: 0.78
--- Confusion Matrix ---
[[2581 | 947]
 [ 775 | 3720]]

```

```

== Training with 31897 examples. For label 'Walking' we have 17735 positive and 14162 negative examples
== Projected the features to 95 features from the sensors: Aud, Acc, Gyro, Loc
== Testing with 8000 examples. For label 'Walking' we have 4401 positive and 3599 negative examples.
-----
Accuracy*:          0.83
Sensitivity (TPR):  0.85
Specificity (TNR):  0.81
Balanced accuracy:  0.83
Precision**:        0.85
F-Score:            0.83
--- Confusion Matrix ---
[[2914 | 685]
 [ 648 | 3753]]

```

5.1 Binary MLPClassifier (Audio Only)

5.2 Binary MLPClassifier (Audio, Accelerometer, Gyroscope, GPS)

6 Discussion

the results achieved demonstrate that for the given model in the given data a 5% increase to the F1 score including the additional data points of accelerometer data gyroscope data and GPS data what it doesn't highlight is how optimal that data is nor does it evaluate the preprocessing of that data nor does it evaluate the generalisation of this methodology to new data sets what it does answer is the original question of does including supplementary data of mu staff increase the F1 score and the answer is yes what is important is to discuss how you would evaluate the generalisation of this model you would also discuss what impact the model quality has on this result if we've used a different model which we have and we will show this result what is the difference in the improvement is there any improvement how do you evaluate whether a model will benefit from this extra data points

Table 8: ExtraSensory Data Breakdown [Vaizman et al. \(2017\)](#)

kfold	1	2	3	4	5
Accuracy	0.92	0.93	0.95	0.94	0.95
Sensitivity (TPR)	0.26	0.35	0.33	0.29	0.38
F-Score	0.9	0.93	0.94	0.93	0.95

there are some models specifically which you do need to do feature reduction and feature analysis on to ensure you don't over specify the features even if you low weight the features you may still have an issue of having too much data and this is just in terms of a performance another element of this is what we got the results what was the cost so we have increase the amount of data we need to use increase the data we need to process increase the data we need to collect to train we are also looking at a larger model file size which for a use case of potentially online processing would be very detrimental so what's the results of 5% are based on solid evidence they may not actually be a valid application or a valid progress for achieving a 5% increase for instance how do we only change the data set and retrain revalidated and then re-evaluated this process would we see similar numbers that is the model likely wouldn't produce the exact same if one school so is our 5% related or proportional to the original benefit had we not done hyperparameter searching to optimise our model potentially we would have seen a much larger game in F1 score from the addition of the additional data points that hypothetical would be because if 1 model is already underperforming / are an optimised then the additional data points May introduce that additional optimisation needed similes conclusion in conclusion what we're looking at is a 5% increase in F1 school which is relevant for this approach for this data set because this data set does closely match our target use case this is actually a valid result and this is a worthy increase to the field what this does raise however is how does hell result generalise across other applications because this project had a specific use case it was not within the scope to evaluate whether or not supplementary data would improve audio event classification and detection in other audio sets simile so the number of audio sets the database database is available is limited and the format of that data all the data selection would play a role so whilst we discussed supplementary data to specific data that we use was axela on the data a gyroscope data GPS however it did not attempt to evaluate other supplementary data is for instance the magna mater other accelerometer data video etcetera this data set is also a 1 minutes sample time snap what we did was we recorded across the day and time time-invariant snapshot of our data why that's important is because would time variance significantly affect this outcome that is something probably to put in the discussion section but in conclusion this seems to be a valid strategy and the next next step you would take is to look at trying a larger data set when available and 2 validate this on hardware in a actual environment further so there were other methods to approach the same problem and it should be considered whether or not those other methods would drastically change the app for instant whilst we used a model that was well well chosen for the purpose of a single binary flag to create an ensemble of classifiers conclusion in conclusion what would looking at is a 5% increase in Heflin school which is relevant for visa parts for this. Affect because the status that sounds call claim at at Target use case this is actually a valid result in this is a wuzzy increase to the failed what is thugs raise however is how does have result general on across the applications because it's project hi to specific use case it was not with him the Skype still valley light weather I'm not supplementary Dada would improve audio event classification inspection in other warriors that simile so the number of audio set device. Basis available is limited and the format of that. Orthodontist election would play a role boat well sweet discussed supplementary gotta to specific daughter that we used was XL alarm the Dada gyroscope. GPS however it did not attempt to evaluate all the supplementary. It Francis magnetic element get to work today sometime

and buried shops all out. Why that's important who's because would come variant significantly effects did outcome that is something on

7 Conclusion

This is one of the most important parts of the report. In the conclusion section, you should

- briefly summarise the results,
- reflect on the work presented,
- make recommendations,
- suggest future work or improvements.

The results achieved showed audio as a sole input for determining physical actions to be a viable but also shows room for improvement. The focus was on what data to use, preprocessing and finally classification. Data use settled on Google's AudioSet [FLAGGED FOR REVIEW] for ease of use, where as the IMU data was needed was recorded through a bespoke app created by myself and a colleague. Preprocessing took some domain knowledge learnt through out Mechatronics, and processed (mel power coefficients). Classification was through machine learning (RNN, known topologies etc.). The work presented does demonstrates a contribution to the audio processing space and explores how mechatronics domain knowledge impacts and extends machine learning. Based on the work presented, it's clear that audio data processing will be a major tool and research topic into the future. The niche explored in this FYP, of amalgamating traditional audio processing with physical sensor input looks to be of limited use, to limited applications. More over, without large datasets of audio paired with physical sensor data, the machine learning tools explored in this FYP will not be able to be used. As such, the rise of audio recording equipped equipment should be the basis for further research into this space. Future work from here would be to look at the larger role of data acquisition. This is both because it was one of the larger hurdles I had to face, plus it enables further research into the effectiveness of what a diverse dataset can bring. Industry will reasonably expect to take aspects of data fusion and machine learning. To know what

8 References and Citations

References

, ??? Tech. rep.

, 2017. Results. (n. d.). DCASE2017.

URL <http://www.cs.tut.fi/sgn/arg/dcase2017/challenge/task-sound-event-detection-in-real-life/results>

, 2018. (n. d.).

Bergstra, J., Bengio, Y., 2012. Random search for hyper-parameter optimization. *Journal of Machine Learning Research* 13 (Feb), 281–305.

Brochu, E., Cora, V. M., de Freitas, N., 2010. A tutorial on bayesian optimization of expensive cost functions, with application to active user modeling and hierarchical reinforcement learning.

Dauphin, Y. N., Fan, A., Auli, M., Grangier, D., 2016. Language modeling with gated convolutional networks. *CoRR*.

DCASE., ??? (2018). Schedule.

URL <http://dcase.community/>

for Standardization, I. O., 2015. IEEE 15288: Systems and software engineering – System life cycle processes. International Organization for Standardization, iee 15288 Edition, international Organization for Standardization. (n.d.). Systems and software engineering – System life cycle processes. Retrieved from <https://www.iso.org/standard/63711.html>.

Geron, A., 2017. Hands on Machine Learning with Scikit, Keras & Tensorflow. O'Reilly.

Google, 2019. Sensors motion.

Google Brain, 2019a. Classification: Roc curve and auc — machine learning crash course — google developers.

URL <https://www.tensorflow.org/install/pip>

Google Brain, 2019b. Get started with tensorflow lite.

URL <https://www.tensorflow.org/install/pip>

Google Brain, 2019c. Install tensorflow with pip — tensorflow.

URL https://www.tensorflow.org/lite/guide/get_started#2_convert_the_model_format

Gruber, H., Holzer, M., Ruepp, O., 2007. Sorting the slow way: An analysis of perversely awful randomized sorting algorithms. In: *Lecture Notes in Computer Science*. Springer Berlin Heidelberg, pp. 183–197.

Havard, S., Oct. 2016. Google pixel xl teardown. https://www.electronicproducts.com/Google_Pixel-whatsinside_text-205.aspx.

Heittola, T., 2018. Research: Sound event detection.

Holzer, M., Maurer, J.-T., 2018. Selection via the bogo-method - more on the analysis of perversely awful randomized algorithms.

- Huffman, D. A., Mar. 1954. The synthesis of sequential switching circuits. J. Franklin Institute, vol. 257, no 257 (3), 161–190.
- James, G., Witten, D., Hastie, T., Tibshirani, R., 2017. An Introduction to Statistical Learning. Springer-Verlag GmbH.
URL https://www.ebook.de/de/product/20292548/gareth_james_daniela_witten_trevor_hastie_robert_tibshirani_an_introduction_to_statistical_learning.html
- JiaKai, L., 2018. Mean teacher convolution system for dcase 2018 task.
- LeCun, Y. A., Bottou, L., Orr, G. B., Muller, K.-R., 2012. Efficient backprop. In: Neural networks: Tricks of the trade. Springer, pp. 9–48.
- Limited, B. H. W., 2015. The Systems Thinking Tool Box. [Online].
- Limited, B. H. W., 2018. Systems Engineering: Design and deliver better products and processes. [Online].
- Magdon-Ismael, M., 2000. No free lunch for noise prediction. Neural Computation 12 (3), 547–564.
URL <https://doi.org/10.1162/089976600300015709>
- Maurice Herlihy, N. S., 2012. The Art of Multiprocessor Programming, Revised Reprint. Elsevier Science.
URL https://www.ebook.de/de/product/19187472/maurice_herlihy_nir_shavit_the_art_of_multiprocessor_programming_revised_reprint.html
- Mitchell, T. M., 1997. Machine Learning. McGraw-Hill Education.
URL <https://www.amazon.com/Machine-Learning-Tom-M-Mitchell/dp/0070428077?SubscriptionId=AKIAIOBINVZYXZQZ2U3A&tag=chimbori05-20&linkCode=xm2&camp=2025&creative=165953&creativeASIN=0070428077>
- of Queensland, T. U., 2018. Literature review. Tech. rep., The University of Queensland.
- Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau, D., Brucher, M., Perrot, M., Duchesnay, E., 2019a. Confusion matrix — scikit-learn 0.21.2 documentation.
- Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau, D., Brucher, M., Perrot, M., Duchesnay, E., 2019b. Release history — scikit-learn 0.21.2 documentation.
URL https://scikit-learn.org/stable/whats_new.html
- Sherman, M., Jun. 2013. Bogo-sort is sort of slow.
- Snoek, J., Larochelle, H., Adams, R. P., 2012. Practical bayesian optimization of machine learning algorithms.
- Vaizman, Y., Ellis, K., Lanckriet, G., oct 2017. Recognizing detailed human context in the wild from smartphones and smartwatches. IEEE Pervasive Computing 16 (4), 62–74.

A Full Extracts of Journals

N/A, email correspondence available on request.

B Reflection

about several of the Engineers Australia 16 competencies. [FLAGGED FOR REVIEW] I'll write this from the first person perspective. This project was a whirlwind. Taken over the 2018 S2, 2019 S1 period, including the Christmas Holidays, the project was a learning experience. The outcome I got is barely quantifiable past a common sense guess, but I can take pride in knowing the outcome is quantified and based on the '360 hours' minimum of work to get to that result. I learnt a lot starting with...

B.1 Lessons Learnt

- The Industry - FYP was a squandered opportunity on my part, but almost by my choice in FYP topic. Work was able to provide significant support for implementation, and even the prospect of data collection through engaging with potential customers and stakeholders for different use-cases. Unfortunately, I felt uncomfortable pursuing that due to how heavy the time spent in implementing felt. The lesson learnt is to have not done an Industry FYP without understanding how it'll affect the pathway to the final thesis.
- Implementation drove some technology choices I probably would have been better off without. E.g. I went to a Matlab seminar on how to develop Machine Learning in Matlab. Had I used this, the project could have been more homogenous and focused on the big question around "Acoustic Event Detection/Classification". But to do so means I don't develop the python ML skills needed for 'real world' ML tasks. A lesson learnt is to focus on what is assessable, and to probably have used Matlab to demonstrate I understood the theory, not to learn and demonstrate the actual skills of Industry.
- The FYP is a 10 week full time commitment to meet the 360 hours. I potentially "wasted" 6 weeks just identifying the question I was trying to answer in the report. If I hadn't elected to do a Industry FYP, this is potentially work that could be better shared between the supervisor and myself.
- Smaller Scope, get it done in Part A, expand and write in Part B.

The FYP has made me appreciate Subject Matter Expertise. In MCHA3900, I focused heavily on understanding the physical hardware involved in the Hexapod - Matlab, RPi, i2c, Openservo v3, and the physical bot itself. I felt I knew that system well, but I recognised my limited understanding. In comparison to that, I reflect on *how much* I've learnt during this FYP, and to what depth of detail compared to MCHA3900. I think I already knew that when I read it's 'MCHA3900 - INTRO TO ROBOTICS', but seeing how that was already a good 200 hours I'm surprised just how much more I've had to delve into this.

On that note, I lastly reflect on the skills I've learnt, and the perspective I've gained in how complex the field is, and how much more there is to learn.




Figure 18: A Random Tree Classifier was used to evaluate feature importance, during the design phase of the main classifier. Features were Audio (26 bins), Accelerometer (), Gyroscope (), GPS ()




Figure 19: A Random Tree Classifier was used to evaluate feature importance, during the design phase of the main classifier. Features were Audio (26 bins), Accelerometer (), Gyroscope (), GPS ()

C Dataset Feature Importance

D Trade-Review of "Online vs Offline" processing

E GridSearch Attempts

This shows how a Grid Search was used for hyperparameter optimisation. In particular, this is an example of an iterative process where the grid was manually narrowed and precision increase.

```
parameters={'hidden_layer_sizes': [(145,), (147,), (149,), (151,), (153)],
            'alpha': [0.09, 0.11, 0.13, 0.15, 0.17, 0.19]}
mlpc = MLPClassifier(verbose=False, early_stopping=True, learning_rate='adaptive', max_iter=1000)
clf = GridSearchCV(estimator=mlpc, scoring='f1', param_grid=parameters, n_jobs=-1, verbose=1, cv=3);
clf.fit(X_train, y);
print(clf.score(X_train, y))
print(clf.best_params_)
```

```
Fitting 3 folds for each of 30 candidates, totalling 90 fits
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 42 tasks | elapsed: 7.8min
[Parallel(n_jobs=-1)]: Done 90 out of 90 | elapsed: 14.4min finished
0.8249400479616308
{'alpha': 0.11, 'hidden_layer_sizes': (151,)}
```




Figure 20: A Random Tree Classifier was used to evaluate feature importance, during the design phase of the main classifier. Features were Audio (26 bins), Accelerometer (), Gyroscope (), GPS ()




Figure 21: A Random Tree Classifier was used to evaluate feature importance, during the design phase of the main classifier. Features were Audio (26 bins), Accelerometer (), Gyroscope (), GPS ()

F Table of ExtraSensory Labels

G Writing you're not sure if you want to keep

G.1 Funny Future Work

A suggested option for Future Work could include adapting the Bogosort, which is an increasingly popular area of research [Gruber et al. \(2007\)](#) [Holzer and Maurer \(2018\)](#) [Sherman \(2013\)](#). The suggested method for using the Bogo search would be to iterate over the hyperparameters of the network, then iterate over the neuron weights. If successfully applied, it could be the first instance of the Bogosort being applied in the Machine Learning sphere.

G.2 Development Tools and Technologies

Windows 10, i7-7500U and GeForce 940MX, Nvidia 1080Ti, VS Code, LaTeX, Python, Matlab, sklearn, numpy, ExtraSensory dataset, git.

Visual Studio was used for developing the Python project

G.3 Discussion on practical aspects of IMU and Audio

- I'm not sure if it's practical to try and be recording IMU plus audio
- It's a case of POV; if you can record IMU, you can also record audio from the POV of your user. This theoretically suggests a higher fidelity in classification, especially in determining whether an event was local (POV) or whether it was an external event (someone else)
- whilst IMU is helpful, the level of understanding

Python was chosen for this project due to its ease of use, ease of analysis, ease of learning and its low maintenance. Choosing which dataset was the most difficult was not difficult for me as I have worked heavily with it and its great use cases are pretty straightforward for large data sets. In fact I even wrote a Python for it recently. Although it is still an early, immature project it works quite well for training models and will be the foundation of others as well. There are very few data in the dataset I could not predict the results from. This gives me confidence that the model is accurate from the start and can be used across all datasets and in other languages.

When the dataset was large the results tended to cluster around the 50th and 75th percentile. From this dataset I learned a bit of things. First off, the most recent observations are generally much better aligned and consistent than the average number of observations. And second this isn't a problem for every one of the training data! It seems most of the time we end up with some data which has better variance and the better alignments in the data. If these outl

To summarize what we did : We load and processed the dataset We got familiar with the dataset by plotting some histograms and a correlation heat map of the features We used a deep neural network with three hidden layers each one has 256 nodes. We used a linear activation function on the output layer We trained the model then test it on Kaggle. We also tested two other models Our deep neural network was able to outscore these two models We believe that these two models could beat the deep neural network model if we tweak their hyperparameters.<https://towardsdatascience.com/deep-neural-networks-for-regression-problems-81321897ca33>

G.4 Nearest Neighbour

Nearest Neighbour was chosen as a classification method to test against, and evaluate the other methods. Nearest Neighbour would not be used as a primary method as does not generalise well^[citation needed]¹

ciation:
Vaizman2017a Vaizman, Y., Ellis, K., and Lanckriet, G. "Recognizing Detailed Human Context In-the-Wild from Smartphones and Smartwatches". IEEE Pervasive Computing, vol. 16, no. 4, October-December 2017, pp. 62-74. doi:10.1109/MPRV.2017.3971131 *Cite this paper if you use the ExtraSensory Dataset for any publication!

sounds like dropout is great for making sure accelerometer data is used

Dropout

Unlike these methods mentioned above, dropout in my understanding is tricky but practical. Like the figure shown below, the dropout will randomly mute some neurons in the neural network and

¹<https://scikit-learn.org/stable/modules/neighbors.html>

we therefore have a sparse network which hugely decreases the possibility of overfitting. More importantly, the dropout will make the weights spread over the input features instead of focusing on some features.

The possibility of muting neurons is often set as 0.5 though you can feel free to make it 0.3 or 1.0. When the dropout is 1.0, then you simply don't drop out any neurons. But our experience tells us 0.5 is usually the best choice.

ExtraSensory Data BreakdoAfter finishing the training, it is important to turn off the dropout during development and testing. Otherwise, the prediction of this model is not stable since dropout add uncertainties to it.

Table 9: ExtraSensory Labels [Vaizman et al. \(2017\)](#)

#	Label Description
1	LYING_DOWN
2	SITTING
3	FIX_walking
4	FIX_running
5	BICYCLING
6	SLEEPING
7	LAB_WORK
8	IN_CLASS
9	IN_A_MEETING
10	LOC_main_workplace
11	OR_indoors
12	OR_outside
13	IN_A_CAR
14	ON_A_BUS
15	DRIVE_-_I_M_THE_DRIVER
16	DRIVE_-_I_M_A_PASSENGER
17	LOC_home
18	FIX_restaurant
19	PHONE_IN_POCKET
20	OR_exercise
21	COOKING
22	SHOPPING
23	STROLLING
24	DRINKING__ALCOHOL_
25	BATHING_-_SHOWER
26	CLEANING
27	DOING_LAUNDRY
28	WASHING_DISHES
29	WATCHING_TV
30	SURFING_THE_INTERNET
31	AT_A_PARTY
32	AT_A_BAR
33	LOC_beach
34	SINGING
35	TALKING
36	COMPUTER_WORK
37	EATING
38	TOILET
39	GROOMING
40	DRESSING
41	AT_THE_GYM
42	STAIRS_-_GOING_UP
43	STAIRS_-_GOING_DOWN
44	ELEVATOR
45	OR_standing
46	AT_SCHOOL
47	PHONE_IN_HAND
48	PHONE_IN_BAG
49	PHONE_ON_TABLE
50	WITH_CO-WORKERS
51	WITH_FRIENDS